

exploration_ce

November 23, 2021

```
[1]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: for dirname, _, filenames in os.walk('../..//
      ↳bigquery-geotab-intersection-congestion'):
      for filename in filenames:
          print(os.path.join(dirname, filename))
```

```
../..//bigquery-geotab-intersection-congestion\BigQuery-Dataset-Access.md
../..//bigquery-geotab-intersection-congestion\sample_submission.csv
../..//bigquery-geotab-intersection-congestion\submission_metric_map
../..//bigquery-geotab-intersection-congestion\submission_metric_map.json
../..//bigquery-geotab-intersection-congestion\test.csv
../..//bigquery-geotab-intersection-congestion\train.csv
```

```
[3]: df_train = pd.read_csv('../..//bigquery-geotab-intersection-congestion/train.
      ↳csv')
df_test = pd.read_csv('../..//bigquery-geotab-intersection-congestion/test.csv')
```

```
[4]: df_train.head()
```

```
[4]:   RowId  IntersectionId  Latitude  Longitude  \
0   1921357             0   33.791659  -84.430032
1   1921358             0   33.791659  -84.430032
2   1921359             0   33.791659  -84.430032
3   1921360             0   33.791659  -84.430032
4   1921361             0   33.791659  -84.430032
```

	EntryStreetName	ExitStreetName	EntryHeading	\
0	Marietta Boulevard Northwest	Marietta Boulevard Northwest	NW	
1	Marietta Boulevard Northwest	Marietta Boulevard Northwest	SE	
2	Marietta Boulevard Northwest	Marietta Boulevard Northwest	NW	
3	Marietta Boulevard Northwest	Marietta Boulevard Northwest	SE	
4	Marietta Boulevard Northwest	Marietta Boulevard Northwest	NW	

	ExitHeading	Hour	Weekend	...	TimeFromFirstStop_p40	\
0	NW	0	0	...	0.0	
1	SE	0	0	...	0.0	
2	NW	1	0	...	0.0	
3	SE	1	0	...	0.0	
4	NW	2	0	...	0.0	

	TimeFromFirstStop_p50	TimeFromFirstStop_p60	TimeFromFirstStop_p80	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	DistanceToFirstStop_p20	DistanceToFirstStop_p40	DistanceToFirstStop_p50	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	DistanceToFirstStop_p60	DistanceToFirstStop_p80	City
0	0.0	0.0	Atlanta
1	0.0	0.0	Atlanta
2	0.0	0.0	Atlanta
3	0.0	0.0	Atlanta
4	0.0	0.0	Atlanta

[5 rows x 28 columns]

```
[5]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 856387 entries, 0 to 856386
Data columns (total 28 columns):
#   Column              Non-Null Count  Dtype
---  -
0   RowId               856387 non-null  int64
1   IntersectionId      856387 non-null  int64
2   Latitude            856387 non-null  float64
3   Longitude           856387 non-null  float64
4   EntryStreetName     848239 non-null  object
5   ExitStreetName      850100 non-null  object
6   EntryHeading        856387 non-null  object
7   ExitHeading         856387 non-null  object
8   Hour                856387 non-null  int64
9   Weekend             856387 non-null  int64
10  Month               856387 non-null  int64
```

```

11 Path 856387 non-null object
12 TotalTimeStopped_p20 856387 non-null float64
13 TotalTimeStopped_p40 856387 non-null float64
14 TotalTimeStopped_p50 856387 non-null float64
15 TotalTimeStopped_p60 856387 non-null float64
16 TotalTimeStopped_p80 856387 non-null float64
17 TimeFromFirstStop_p20 856387 non-null float64
18 TimeFromFirstStop_p40 856387 non-null float64
19 TimeFromFirstStop_p50 856387 non-null float64
20 TimeFromFirstStop_p60 856387 non-null float64
21 TimeFromFirstStop_p80 856387 non-null float64
22 DistanceToFirstStop_p20 856387 non-null float64
23 DistanceToFirstStop_p40 856387 non-null float64
24 DistanceToFirstStop_p50 856387 non-null float64
25 DistanceToFirstStop_p60 856387 non-null float64
26 DistanceToFirstStop_p80 856387 non-null float64
27 City 856387 non-null object

```

dtypes: float64(17), int64(5), object(6)

memory usage: 182.9+ MB

```

[6]: df_train.dropna()
df_test.dropna()

summary_table = df_train.describe()
summary_table

```

```

[6]:
count      RowId  IntersectionId      Latitude      Longitude \
count  8.563870e+05  856387.000000  856387.000000  856387.000000
mean    2.349550e+06    833.283384    39.618965   -77.916488
std     2.472178e+05    654.308913     2.935437     5.952959
min     1.921357e+06     0.000000    33.649973   -87.862288
25%     2.135454e+06    291.000000    39.936739   -84.387607
50%     2.349550e+06    679.000000    39.982974   -75.175055
75%     2.563646e+06   1264.000000    41.910047   -75.100495
max     2.777743e+06   2875.000000    42.381782   -71.025550

count      Hour      Weekend      Month  TotalTimeStopped_p20 \
count  856387.000000  856387.000000  856387.000000  856387.000000
mean    12.431234    0.277880     9.104808     1.755596
std      6.071843    0.447954     1.991094     7.146549
min      0.000000    0.000000     1.000000     0.000000
25%      8.000000    0.000000     7.000000     0.000000
50%     13.000000    0.000000     9.000000     0.000000
75%     17.000000    1.000000    11.000000     0.000000
max     23.000000    1.000000    12.000000    298.000000

count      TotalTimeStopped_p40  TotalTimeStopped_p50  ...  TimeFromFirstStop_p20 \
count      856387.000000      856387.000000  ...      856387.000000

```

mean	5.403592	7.722655	...	3.181096
std	12.981674	15.685910	...	11.835994
min	0.000000	0.000000	...	0.000000
25%	0.000000	0.000000	...	0.000000
50%	0.000000	0.000000	...	0.000000
75%	0.000000	10.000000	...	0.000000
max	375.000000	375.000000	...	337.000000

	TimeFromFirstStop_p40	TimeFromFirstStop_p50	TimeFromFirstStop_p60	\
count	856387.000000	856387.000000	856387.000000	
mean	9.162174	12.722165	18.926085	
std	20.446568	24.219271	29.851797	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	
75%	0.000000	22.000000	31.000000	
max	356.000000	356.000000	357.000000	

	TimeFromFirstStop_p80	DistanceToFirstStop_p20	\
count	856387.000000	856387.000000	
mean	34.201656	6.765856	
std	41.130668	29.535968	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	27.000000	0.000000	
75%	49.000000	0.000000	
max	359.000000	1901.900000	

	DistanceToFirstStop_p40	DistanceToFirstStop_p50	\
count	856387.000000	856387.000000	
mean	20.285128	28.837113	
std	59.202108	75.217343	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	0.000000	53.100000	
max	2844.400000	2851.100000	


	DistanceToFirstStop_p60	DistanceToFirstStop_p80
count	856387.000000	856387.000000
mean	44.27231	83.991313
std	102.03225	160.709797
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	60.400000
75%	64.20000	85.950000
max	3282.40000	4079.200000

[8 rows x 22 columns]

```
[7]: names = ['TotalTimeStopped_p20',
              'TotalTimeStopped_p40',
              'TotalTimeStopped_p50',
              'TotalTimeStopped_p60',
              'TotalTimeStopped_p80']
x = [summary_table[name]['mean'] for name in names]
labels = [name[-3:] for name in names]

plt.barh(labels, x)

plt.show()
```



```
[8]: numerical_cols = [col for col in summary_table.columns]
categorical_cols = [col for col in df_train.columns if col not in
                    numerical_cols]

for col in categorical_cols:
    print(col)
```

EntryStreetName
ExitStreetName
EntryHeading
ExitHeading

Path
City

```
[9]: # Dummy variables

# Training
df_train = pd.concat([df_train, pd.get_dummies(df_train['EntryHeading'],
                                              prefix='entry')],
                    axis=1)
df_train = pd.concat([df_train, pd.get_dummies(df_train['ExitHeading'],
                                              prefix='exit')],
                    axis=1)
df_train = pd.concat([df_train, pd.get_dummies(df_train['City'],
                                              prefix='c')],
                    axis=1)

# Test
df_test = pd.concat([df_test, pd.get_dummies(df_test['EntryHeading'],
                                              prefix='entry')],
                    axis=1)
df_test = pd.concat([df_test, pd.get_dummies(df_test['ExitHeading'],
                                              prefix='exit')],
                    axis=1)
df_test = pd.concat([df_test, pd.get_dummies(df_test['City'],
                                              prefix='c')],
                    axis=1)
```

```
[10]: numerical_cols = [col for col in summary_table.columns]
categorical_cols = [col for col in df_train.columns if col not in_
    ↪ numerical_cols]

for col in categorical_cols:
    print(col)
```

EntryStreetName
ExitStreetName
EntryHeading
ExitHeading
Path
City
entry_E
entry_N
entry_NE
entry_NW
entry_S
entry_SE
entry_SW
entry_W
exit_E

```

exit_N
exit_NE
exit_NW
exit_S
exit_SE
exit_SW
exit_W
c_Atlanta
c_Boston
c_Chicago
c_Philadelphia

```

```


[11]: targets = ['TotalTimeStopped', 'TimeFromFirstStop', 'DistanceToFirstStop']
percentiles = ['p20', 'p40', 'p50', 'p60', 'p80']

plt.figure(figsize=(16, 4))
kwargs = dict(alpha=0.5, bins=100)
xlims = [200, 300, 1500]

i = 1
for target in targets:
    plt.subplot(int('13' + str(i)), label=target)
    for percentile in percentiles:
        plt.hist(df_train[target+'_'+percentile], **kwargs, label=percentile)
    plt.gca().set(title=target, ylabel='Frequency')
    plt.xlim(0, xlims[i-1])
    plt.legend()
    i += 1

plt.subplots_adjust(right = 1)
plt.show()

```



```

[12]: from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.multioutput import MultiOutputRegressor as MOR

from sklearn.ensemble import GradientBoostingRegressor

```

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression

```

```

[13]: X = df_train[["IntersectionId", "Hour", "Weekend", "Month", 'entry_E', 'entry_N',
    ↳ 'entry_NE',
    ↳ 'entry_NW', 'entry_S', 'entry_SE', 'entry_SW', 'entry_W',
    ↳ 'exit_E', 'exit_N',
    ↳ 'exit_NE', 'exit_NW', 'exit_S', 'exit_SE', 'exit_SW', 'exit_W',
    ↳ 'c_Atlanta', 'c_Boston', 'c_Chicago', 'c_Philadelphia']]
y = pd.concat([df_train[f'{targets[0]}_{perc}'] for perc in percentiles],
    ↳ axis=1)

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=.2)

X_test = df_test[["IntersectionId", "Hour", "Weekend", "Month", 'entry_E',
    ↳ 'entry_N', 'entry_NE',
    ↳ 'entry_NW', 'entry_S', 'entry_SE', 'entry_SW', 'entry_W',
    ↳ 'exit_E', 'exit_N',
    ↳ 'exit_NE', 'exit_NW', 'exit_S', 'exit_SE', 'exit_SW',
    ↳ 'exit_W',
    ↳ 'c_Atlanta', 'c_Boston', 'c_Chicago', 'c_Philadelphia']]

```

```

[15]: # TODO research good regressors for this data
models = {
    'gbr': GradientBoostingRegressor,
    'rfr': RandomForestRegressor,
    'lr': LinearRegression
}
# TODO research hyperparameters for the models
# params = {
#     'gbr': {
#         'loss': ['squared_error', 'absolute_error', 'huber', 'quantile'],
#         'learning_rate': [.01, .1,]
#     }
# }
# }
# TODO: Normalization, pipelining

```

```

[16]: preds = {}
for model in models:
    preds[model] = np.asarray(MOR(models[model]()).fit(X_train, y_train).
    ↳ predict(X_val))

```

```

[17]: for model in models:
    print(model)
    i = 0
    for col in y_val:
        print(metrics.mean_squared_error(y_val[col].values, preds[model].T[i],
    ↳ squared=False))

```



```
i += 1  
print()
```

gbr

7.054760286217763
12.579387954820882
15.004284745473528
18.78764257177234
26.54303354817789

rfr

6.135557352863927
9.365125300660246
10.639747714059608
12.829687150920739
18.07481494270436

lr

7.200918975835944
13.036677664109
15.6153572658614
19.594620751784458
27.776653280428995

```
[ ]: # I set up test but never used it
```